

```
In [1]: import pandas as pd
import numpy as np
import csv
import random as rd
```

```
In [2]: df_movies = pd.read_csv("C:\\Users\\amang\\Downloads\\task2\\movies1.csv", sep=
df_movies.columns = ['Movie_IDs', 'Movie_Name', 'Genre']
df_movies.dropna(inplace=True)
df_movies = df_movies.replace('','', regex=True)
df_movies.head()
```

Out[2]:

	Movie_IDs	Movie_Name	Genre
0	2	Jumanji (1995)	Adventure Children's Fantasy
1	3	Grumpier Old Men (1995)	Comedy Romance
2	4	Waiting to Exhale (1995)	Comedy Drama
3	5	Father of the Bride Part II (1995)	Comedy
4	6	Heat (1995)	Action Crime Thriller

```
In [3]: df_rating = pd.read_csv("C:\\Users\\amang\\Downloads\\task2\\ratings1.csv", sep=
df_rating.columns = ['ID', 'Movies_ID', 'Rating', 'Timestamp']
df_rating.dropna(inplace=True)
df_rating = df_rating.replace('','', regex=True)
df_rating.head()
```

Out[3]:

	ID	Movies_ID	Rating	Timestamp
0	1	661	3	978302109
1	1	914	3	978301968
2	1	3408	4	978300275
3	1	2355	5	978824291
4	1	1197	3	978302268

```
In [4]: df_user = pd.read_csv("C:\\Users\\amang\\Downloads\\task2\\users1.csv", sep='::')
df_user.columns = ['UserID', 'Gender', 'Age', 'Occupation', 'Zip-code']
df_user.dropna(inplace=True)
df_user = df_user.replace(' ', '', regex=True)
df_user.head()
```

Out[4]:

	UserID	Gender	Age	Occupation	Zip-code
0	2	M	56	16	70072
1	3	M	25	15	55117
2	4	M	45	7	02460
3	5	M	25	20	55455
4	6	F	50	9	55117

```
In [5]: data = pd.concat([df_movies, df_rating, df_user], axis = 1)
data.head()
```

Out[5]:

	Movie_IDs	Movie_Name	Genre	ID	Movies_ID	Rating	Timestamp	UserID
0	2	Jumanji (1995)	Adventure Children's Fantasy	1	661	3	978302109	2
1	3	Grumpier Old Men (1995)	Comedy Romance	1	914	3	978301968	3
2	4	Waiting to Exhale (1995)	Comedy Drama	1	3408	4	978300275	4
3	5	Father of the Bride Part II (1995)	Comedy	1	2355	5	978824291	5
4	6	Heat (1995)	Action Crime Thriller	1	1197	3	978302268	6

In [6]: data.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1000208 entries, 0 to 1000207
Data columns (total 12 columns):
 #   Column          Non-Null Count  Dtype
---  --
 0   Movie_IDs       3832 non-null   object
 1   Movie_Name      3832 non-null   object
 2   Genre           3832 non-null   object
 3   ID              1000208 non-null object
 4   Movies_ID       1000208 non-null int64
 5   Rating          1000208 non-null int64
 6   Timestamp       1000208 non-null object
 7   UserID          6039 non-null   object
 8   Gender          6039 non-null   object
 9   Age            6039 non-null   float64
10  Occupation      6039 non-null   float64
11  Zip-code        6039 non-null   object
dtypes: float64(2), int64(2), object(8)
memory usage: 99.2+ MB
```

In [7]: data.describe(include = "all")

Out[7]:

	Movie_IDs	Movie_Name	Genre	ID	Movies_ID	Rating	Timestamp	UserI
count	3832	3832	3832	1000208	1.000208e+06	1.000208e+06	1000208	603
unique	3832	3832	301	6040	NaN	NaN	458455	603
top	2	Jumanji (1995)	Drama	4169	NaN	NaN	975528402	
freq	1	1	822	2314	NaN	NaN	30	
mean	NaN	NaN	NaN	NaN	1.865541e+03	3.581563e+00	NaN	Na
std	NaN	NaN	NaN	NaN	1.096041e+03	1.117102e+00	NaN	Na
min	NaN	NaN	NaN	NaN	1.000000e+00	1.000000e+00	NaN	Na
25%	NaN	NaN	NaN	NaN	1.030000e+03	3.000000e+00	NaN	Na
50%	NaN	NaN	NaN	NaN	1.835000e+03	4.000000e+00	NaN	Na
75%	NaN	NaN	NaN	NaN	2.770000e+03	4.000000e+00	NaN	Na
max	NaN	NaN	NaN	NaN	3.952000e+03	5.000000e+00	NaN	Na

In [8]: `data.isna()`

Out[8]:

	Movie_IDs	Movie_Name	Genre	ID	Movies_ID	Rating	Timestamp	UserID	Gender
0	False	False	False	False	False	False	False	False	False
1	False	False	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False	False	False
3	False	False	False	False	False	False	False	False	False
4	False	False	False	False	False	False	False	False	False
...
1000203	True	True	True	False	False	False	False	True	True
1000204	True	True	True	False	False	False	False	True	True
1000205	True	True	True	False	False	False	False	True	True
1000206	True	True	True	False	False	False	False	True	True
1000207	True	True	True	False	False	False	False	True	True

1000208 rows × 12 columns



In [9]: `data.isna().sum()`

Out[9]:


Movie_IDs	996376
Movie_Name	996376
Genre	996376
ID	0
Movies_ID	0
Rating	0
Timestamp	0
UserID	994169
Gender	994169
Age	994169
Occupation	994169
Zip-code	994169

dtype: int64

```
In [10]: data.dropna(axis=0,inplace=True)
data.isna().sum()
```

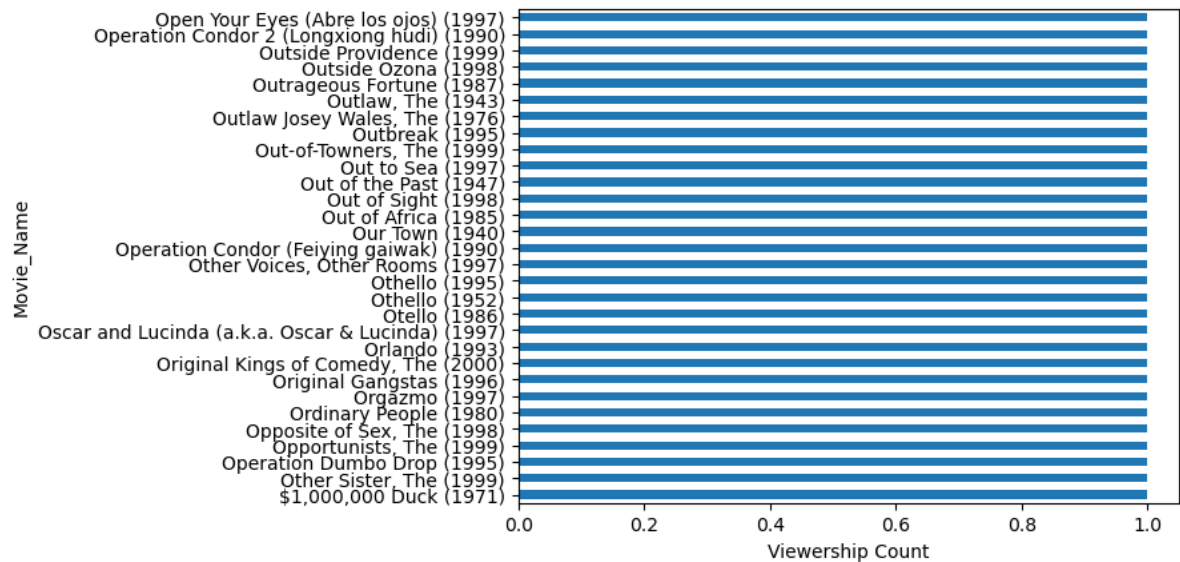
```
Out[10]: Movie_IDs      0
Movie_Name    0
Genre         0
ID            0
Movies_ID     0
Rating        0
Timestamp     0
UserID        0
Gender        0
Age           0
Occupation    0
Zip-code      0
dtype: int64
```

```
In [11]: import matplotlib.pyplot as plt
df_user['Age'].hist(bins=25)
plt.xlabel('Age')
plt.ylabel('Population')
plt.show()
```



```
In [12]: res = data.groupby("Movie_Name").size().sort_values(ascending=False)[:30]
plt.ylabel("Title")
plt.xlabel("Viewership Count")
res.plot(kind="barh")
```

Out[12]: <AxesSubplot:xlabel='Viewership Count', ylabel='Movie_Name'>



```
In [13]: data['Rating'].unique()
```

Out[13]: array([3, 4, 5, 2, 1], dtype=int64)

```
In [14]: df = data['Genre'].str.get_dummies(sep='|')
df.head()
```

Out[14]:

	Action	Adventure	Animation	Children's	Comedy	Crime	Documentary	Drama	Fantasy	Film No.
0	0	1	0	1	0	0	0	0	1	
1	0	0	0	0	1	0	0	0	0	
2	0	0	0	0	1	0	0	1	0	
3	0	0	0	0	1	0	0	0	0	
4	1	0	0	0	0	1	0	0	0	

```
In [15]: df = pd.concat((df,data['Rating']),axis=1)
df.head()
```

Out[15]:

	Action	Adventure	Animation	Children's	Comedy	Crime	Documentary	Drama	Fantasy	File No
0	0	1	0	1	0	0	0	0	1	
1	0	0	0	0	1	0	0	0	0	
2	0	0	0	0	1	0	0	1	0	
3	0	0	0	0	1	0	0	0	0	
4	1	0	0	0	0	1	0	0	0	

```
In [43]: df = pd.concat((df,data['Gender']),axis = 1)
df = pd.concat((df,data['Age']),axis = 1)
df.head()
df = df.loc[:,~df.columns.duplicated()]
df
```

Out[43]:

	Action	Adventure	Animation	Children's	Comedy	Crime	Documentary	Drama	Fantasy
0	0	1	0	1	0	0	0	0	1
1	0	0	0	0	1	0	0	0	0
2	0	0	0	0	1	0	0	1	0
3	0	0	0	0	1	0	0	0	0
4	1	0	0	0	0	1	0	0	0
...
3877	0	0	0	0	1	0	0	0	0
3878	0	0	0	0	0	0	0	1	0
3879	0	0	0	0	0	0	0	1	0
3880	0	0	0	0	0	0	0	1	0
3881	0	0	0	0	0	0	0	1	0

3832 rows × 21 columns

```
In [17]: from sklearn.model_selection import train_test_split
from sklearn.impute import SimpleImputer
from sklearn.metrics import accuracy_score
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error
```

C:\ProgramData\Anaconda3\lib\site-packages\scipy__init__.py:155: UserWarning: A NumPy version >=1.18.5 and <1.25.0 is required for this version of SciPy (detected version 1.25.0
warnings.warn(f"A NumPy version >={np_minversion} and <{np_maxversion}")

```
In [18]: X = df.drop(["Rating", "Age", "Gender"], axis=1)
y = df["Rating"]
```

```
In [34]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.4, random_state=42)
```

```
In [35]: model = LinearRegression()
```

```
In [36]: model.fit(X_train, y_train)
```

```
Out[36]: LinearRegression()
```

```
In [37]: y_pred = model.predict(X_test)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
print(mse)
print(rmse)
```

1.2071128036499075
1.098686854226402

#Decision tree

```
In [38]: from sklearn.tree import DecisionTreeRegressor
model = DecisionTreeRegressor()
model.fit(X_train, y_train)
```

```
Out[38]: DecisionTreeRegressor()
```



```
In [39]: y_pred = model.predict(X_test)

mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f"Mean Absolute Error: {mae}")
print(f"Mean Squared Error: {mse}")
print(f"R-squared: {r2}")
```

```
Mean Absolute Error: 0.939357165603892
Mean Squared Error: 1.3132902142011524
R-squared: -0.09512489832448923
```